

# Foundations of Data Science & Analytics: Ensemble Learning

Ezgi Siir Kibris

[Introduction to Data Mining, 2nd Edition](#)

by

Tan, Steinbach, Karpatne, Kumar

# Classification Techniques

- **Base Classifiers**

- Decision Tree based Methods
- Rule-based Methods
- Instance-based Methods (Nearest-neighbor)
- Naïve Bayes
- Support Vector Machines
- Neural Networks and Deep Learning

- **Ensemble Classifiers**

- **Boosting, Bagging, Random Forests**


# Ensemble Learning

- Construct a set of (weak) classifiers from the training data
- Predict class label of test records by combining the predictions made by multiple classifiers

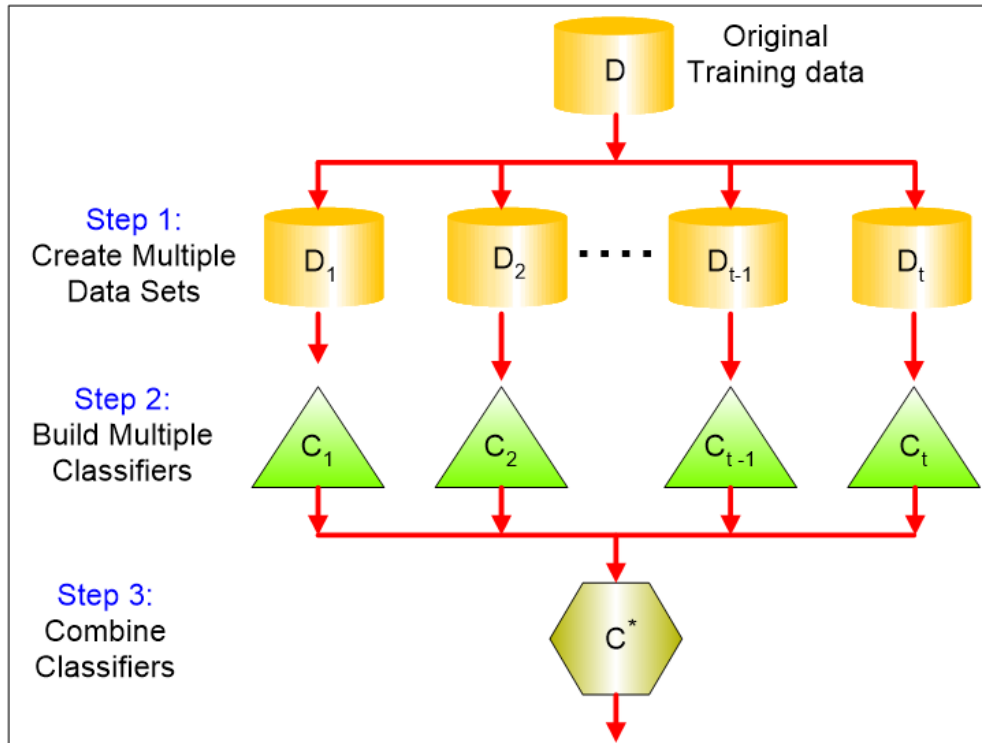
# Why?

- Suppose there are 25 base classifiers
  - Each classifier has error rate,  $\epsilon = 0.35$
  - Assume errors made by classifiers are uncorrelated
  - Probability that the ensemble classifier makes a wrong prediction:

$$P(X \geq 13) = \sum_{i=13}^{25} \binom{25}{i} \epsilon^i (1 - \epsilon)^{25-i} = 0.06$$



# General Approach



# Types

- Manipulate data distribution
  - Bagging
  - Boosting
- Manipulate input features
  - Random Forests

# Types

- Manipulate data distribution
  - **Bagging**
  - Boosting
- Manipulate input features
  - Random Forests

# Bagging (Bootstrap Aggregation)

- Sample with replacement ( $n = N$ )

Original Data	1	2	3	4	5	6	7	8	9	10
Bagging (Round 1)	7	8	10	8	2	5	10	10	5	9
Bagging (Round 2)	1	4	9	1	2	3	2	7	3	2
Bagging (Round 3)	1	8	5	10	5	5	9	6	3	7

- Build classifier on each bootstrap sample
- Each data instance has probability  $1/n$  of being selected each time
- Each data instance has probability  $1 - (1 - 1/n)^n$  of being selected as part of the bootstrap sample



# Bagging

---

## Algorithm 5.6 Bagging Algorithm

---

- 1: Let  $k$  be the number of bootstrap samples.
  - 2: for  $i = 1$  to  $k$  do
  - 3:   Create a bootstrap sample of size  $n$ ,  $D_i$ .
  - 4:   Train a base classifier  $C_i$  on the bootstrap sample  $D_i$ .
  - 5: end for
  - 6:  $C^*(x) = \arg \max_y \sum_i \delta(C_i(x) = y)$ ,  $\{\delta(\cdot) = 1$  if its argument is true, and 0 otherwise. $\}$
- 


Each classifier can be trained in parallel!

# Bagging Example

Original Data:

x	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1
y	1	1	1	-1	-1	-1	-1	1	1	1

- Classifier is a decision stump
  - Decision rule:  $x \leq k$  vs  $x > k$
  - Split point  $k$  is chosen based on entropy



# Bagging Example

Bagging Round 1:

x	0.1	0.2	0.2	0.3	0.4	0.4	0.5	0.6	0.9	0.9
y	1	1	1	1	-1	-1	-1	-1	1	1

$x \leq 0.35 \rightarrow y = 1$

$x > 0.35 \rightarrow y = -1$

# Bagging Example

Bagging Round 1:

x	0.1	0.2	0.2	0.3	0.4	0.4	0.5	0.6	0.9	0.9
y	1	1	1	1	-1	-1	-1	-1	1	1

 $x \leq 0.35 \rightarrow y = 1$  $x > 0.35 \rightarrow y = -1$ 

Bagging Round 2:

x	0.1	0.2	0.3	0.4	0.5	0.5	0.9	1	1	1
y	1	1	1	-1	-1	-1	1	1	1	1

 $x \leq 0.7 \rightarrow y = 1$  $x > 0.7 \rightarrow y = -1$ 

Bagging Round 3:

x	0.1	0.2	0.3	0.4	0.4	0.5	0.7	0.7	0.8	0.9
y	1	1	1	-1	-1	-1	-1	-1	1	1

 $x \leq 0.35 \rightarrow y = 1$  $x > 0.35 \rightarrow y = -1$ 

Bagging Round 4:

x	0.1	0.1	0.2	0.4	0.4	0.5	0.5	0.7	0.8	0.9
y	1	1	1	-1	-1	-1	-1	-1	1	1

 $x \leq 0.3 \rightarrow y = 1$  $x > 0.3 \rightarrow y = -1$ 

Bagging Round 5:

x	0.1	0.1	0.2	0.5	0.6	0.6	0.6	1	1	1
y	1	1	1	-1	-1	-1	-1	1	1	1

 $x \leq 0.35 \rightarrow y = 1$  $x > 0.35 \rightarrow y = -1$

# Bagging Example

Bagging Round 6:

x	0.2	0.4	0.5	0.6	0.7	0.7	0.7	0.8	0.9	1
y	1	-1	-1	-1	-1	-1	-1	1	1	1

 $x \leq 0.75 \rightarrow y = -1$  $x > 0.75 \rightarrow y = 1$ 

Bagging Round 7:

x	0.1	0.4	0.4	0.6	0.7	0.8	0.9	0.9	0.9	1
y	1	-1	-1	-1	-1	1	1	1	1	1

 $x \leq 0.75 \rightarrow y = -1$  $x > 0.75 \rightarrow y = 1$ 

Bagging Round 8:

x	0.1	0.2	0.5	0.5	0.5	0.7	0.7	0.8	0.9	1
y	1	1	-1	-1	-1	-1	-1	1	1	1

 $x \leq 0.75 \rightarrow y = -1$  $x > 0.75 \rightarrow y = 1$ 

Bagging Round 9:

x	0.1	0.3	0.4	0.4	0.6	0.7	0.7	0.8	1	1
y	1	1	-1	-1	-1	-1	-1	1	1	1

 $x \leq 0.75 \rightarrow y = -1$  $x > 0.75 \rightarrow y = 1$ 

Bagging Round 10:

x	0.1	0.1	0.1	0.1	0.3	0.3	0.8	0.8	0.9	0.9
y	1	1	1	1	1	1	1	1	1	1

 $x \leq 0.05 \rightarrow y = 1$  $x > 0.05 \rightarrow y = 1$

# Test on Training Data (with Majority Vote)

Original Data:

x	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1
y	1	1	1	-1	-1	-1	-1	1	1	1

Round	Split Point	Left Class	Right Class
1	0.35	1	-1
2	0.7	1	1
3	0.35	1	-1
4	0.3	1	-1
5	0.35	1	-1
6	0.75	-1	1
7	0.75	-1	1
8	0.75	-1	1
9	0.75	-1	1
10	0.05	1	1

# Types

- Manipulate data distribution
  - Bagging
  - **Boosting**
- Manipulate input features
  - Random Forests

# Boosting

- **Weights:** probability of being sampled
- Records that are wrongly classified will have their **weights** increased
- Records that are classified correctly will have their **weights** decreased

Original Data	1	2	3	4	5	6	7	8	9	10
Boosting (Round 1)	7	3	2	8	7	9	4	10	6	3
Boosting (Round 2)	5	4	9	4	2	5	1	7	4	2
Boosting (Round 3)	4	4	8	10	4	5	4	6	3	4

- Example 4 is hard to classify
- Its weight is increased, therefore it is more likely to be chosen again in subsequent rounds



# AdaBoost

- If any intermediate rounds produce error rate higher than 50%, the weights are reverted back to  $1/n$  and the resampling procedure is repeated (otherwise there will be  $\alpha_i < 0$ )
- Classification: 
$$C^*(x) = \operatorname{argmax}_y \sum_{i=1}^T \alpha_i \delta(C_i(x) = y)$$

# AdaBoost

---

**Algorithm 5.7** AdaBoost Algorithm

---


- 1:  $w = \{w_j = 1/n \mid j = 1, 2, \dots, n\}$ .    {Initialize the weights for all  $n$  instances.}
  - 2: Let  $k$  be the number of boosting rounds.
  - 3: for  $i = 1$  to  $k$  do
  - 4:    Create training set  $D_i$  by sampling (with replacement) from  $D$  according to  $w$ .
  - 5:    Train a base classifier  $C_i$  on  $D_i$ .
  - 6:    Apply  $C_i$  to all instances in the original training set,  $D$ .
  - 7:     $\epsilon_i = \left[ \sum_j w_j \delta(C_i(x_j) \neq y_j) \right]$     {Calculate the weighted error}
  - 8:    if  $\epsilon_i > 0.5$  then
  - 9:      $w = \{w_j = 1/n \mid j = 1, 2, \dots, n\}$ .    {Reset the weights for all  $n$  instances.}
  - 10:    Go back to Step 4.
  - 11:    end if
  - 12:     $\alpha_i = \ln \frac{1-\epsilon_i}{\epsilon_i}$ .
  - 13:    Update the weight of each instance according to equation (5.88).
  - 14: end for
  - 15:  $C^*(x) = \arg \max_y \sum_{j=1}^T \alpha_j \delta(C_j(x) = y)$ .
-

# AdaBoost Example

Original Data:

x	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1
y	1	1	1	-1	-1	-1	-1	1	1	1

- Classifier is a decision stump
  - Decision rule:  $x \leq k$  vs  $x > k$
  - Split point  $k$  is chosen based on entropy



# AdaBoost Example

Original Data:

x	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1
y	1	1	1	-1	-1	-1	-1	1	1	1

Boosting Round 1:

x	0.1	0.4	0.5	0.6	0.6	0.7	0.7	0.7	0.8	1
y	1	-1	-1	-1	-1	-1	-1	-1	1	1

Training:

Weights:

Round	x=0.1	x=0.2	x=0.3	x=0.4	x=0.5	x=0.6	x=0.7	x=0.8	x=0.9	x=1.0
1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1
2	0.167	0.167	0.167	0.071	0.071	0.071	0.071	0.071	0.071	0.071
3	0.117	0.117	0.117	0.125	0.125	0.125	0.125	0.050	0.050	0.050

# AdaBoost Example

Model:

Round	Split Point	Left Class	Right Class	alpha
1	0.75	-1	1	0.847
2	0.05	1	1	0.924
3	0.3	1	-1	1.735

Predictions:

Round	x=0.1	x=0.2	x=0.3	x=0.4	x=0.5	x=0.6	x=0.7	x=0.8	x=0.9	x=1.0
1	-1	-1	-1	-1	-1	-1	-1	1	1	1
2	1	1	1	1	1	1	1	1	1	1
3	1	1	1	-1	-1	-1	-1	-1	-1	-1
Sum	1.81	1.81	1.81	-1.7	-1.7	-1.7	-1.7	0.04	0.04	0.04
Sign	1	1	1	-1	-1	-1	-1	1	1	1

# Further Read on Adaboost

<https://web.stanford.edu/~hastie/Papers/samme.pdf>

# Types

- Manipulate data distribution
  - Bagging
  - Boosting
- Manipulate input features
  - **Random Forests**

# Random Forest

Each tree is trained on a subset of data (bagging) with a subset of features (feature bagging).

- Bagging (sample with replacement)
  - $n = N$
- Feature bagging (sample with replacement)
  - usually with  $m = \sqrt{M}$  features.



# Random Forest


1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Bagging



13	14	15	16
1	2	3	4
1	2	3	4
5	6	7	8

Feature  
Bagging



14	15
2	3
2	3
6	7